

# Machine Learning Canvas

	PREDICTIONS	OBJECTIVES	DATA
IDEA	<p><b>Context</b></p> <p>Who will use the predictive system / who will be affected by it? Provide some background.</p> <p>Airbnb hosts and platform admins.</p> <p>Hosts need to set optimal rental prices to maximize occupancy and revenue. Platform admins can benefit from insights into pricing trends and patterns to improve service offerings.</p>	<p><b>Value Proposition</b></p> <p>What are we trying to do? E.g. spend less time on X, increase Y...</p> <p>Optimize rental pricing to increase occupancy rates and revenue for Airbnb hosts.</p>	<p><b>Data Sources</b></p> <p>Where do/can we get data from? (internal database, 3rd party API, etc.)</p> <p>Kaggle dataset containing Airbnb listings data.</p>
	<p><b>Problem</b></p> <p>Question to predict answers to (in plain English)</p> <p>How much should an Airbnb host charge for a rental property to maximize revenue and occupancy?</p> <p>Input (i.e. question "parameter")</p> <p>Property features; Time-related data; Market conditions.</p> <p>Possible outputs (i.e. "answers")</p> <p>Optimal rental price for a given property at a specific time.</p> <p>Type of problem (e.g. classification, regression, recommendation...)</p> <p>Regression problem (predicting a continuous variable: price).</p>	<p><b>Performance evaluation</b></p> <p>Domain-specific / bottom-line metrics for monitoring performance in production</p> <p>Increase in average daily revenue, occupancy rate improvement.</p> <p>Prediction accuracy metrics (e.g. MSE if regression; % accuracy, #FP for classification)</p> <p>Mean Squared Error (MSE): Measures average squared difference between predicted and actual prices.</p> <p>Mean Absolute Error (MAE): Measures average absolute difference between predicted and actual prices.</p> <p>Offline performance evaluation method (e.g. cross-validation or simple training/test split)</p> <p>Cross-Validation: Ensures model robustness and generalizability.</p> <p>Train/Test Split: Provides initial performance benchmarks.</p> <p>Dataset</p>	<p><b>Dataset</b></p> <p>How do we collect data (inputs and outputs)? How many data points?</p> <p>Extract data from Kaggle dataset, preprocess and clean.</p> <p>Initial dataset contains 74,111 rows with 29 columns of various features.</p> <p><b>Features</b></p> <p>Used to represent inputs and extracted from data sources above. Group by types and mention key features if too many to list all.</p> <p>Numerical: Price, number of rooms, number of beds, etc.</p> <p>Categorical: Property type, room type, neighborhood.</p> <p>Text: Descriptions and reviews</p> <p>Datetime: Booking date, listing date, etc.</p> <p>Key Features: Location, number of rooms, property type, amenities, season, reviews.</p>
SPECS	<p><b>Baseline</b></p> <p>What is an alternative way of making predictions (e.g. manual rules based on feature values)?</p> <p>Hosts use personal experience and historical data to set prices.</p>		
DEPLOYMENT	<p><b>Using predictions</b></p> <p>When do we make predictions and how many?</p> <p>Predictions are made whenever a host lists or updates a property.</p> <p>What is the time constraint for making those predictions?</p> <p>Predictions should be quick to ensure a seamless user experience.</p> <p>How do we use predictions and confidence values?</p> <p>Provide suggested prices with confidence intervals to hosts for final decision-making.</p>	<p><b>Learning predictive models</b></p> <p>When do we create/update models? With which data / how much?</p> <p>Models should be updated periodically (e.g., monthly) or when significant new data is available. The most recent data is used.</p> <p>What is the time constraint for creating a model?</p> <p>Few hours ideally.</p> <p>Criteria for deploying model (e.g. minimum performance value — absolute, relative to baseline or to previous model)</p> <p>Model must meet or exceed our success criteria metrics. Model must perform consistently across different time periods and property types.</p>	

Reset Form

Machine Learning Canvas v0.1

[Louis Dorard](#) © 2015. All rights reserved. Please reference [machinelearningcanvas.com](http://machinelearningcanvas.com) by linking to it if you use the canvas.